

# Lab 0: Design Doc

## Preliminaries

## Booting Pintos

## Debugging

## QUESTIONS: BIOS

## QUESTIONS: BOOTLOADER

## QUESTIONS: KERNEL

## Kernel Monitor

## Preliminaries

- Name: 张龄心
- Email: 2100013018@stu.pku.edu.cn
- Reference:

[https://blog.csdn.net/u011011827/article/details/125656962?ops\\_request\\_misc={\"request\\_id\":\"167772676616800188575949\",\"scm\":\"20140713.130102334.pc\\_all.\"}&request\\_id=167772676616800188575949&task-code-2-all-first\\_rank\\_ecpm\\_v1~rank\\_v31\\_ecpm-1-125656962-0-null-null.142^v73^insert\\_down1,201^v4^add\\_ask,239^v2^insert\\_chatgpt&utm\\_term=shell中实现backspace](https://blog.csdn.net/u011011827/article/details/125656962?ops_request_misc={\)  
[https://en.wikipedia.org/wiki/BIOS\\_interrupt\\_call#/Interrupt\\_table](https://en.wikipedia.org/wiki/BIOS_interrupt_call#/Interrupt_table)

## Booting Pintos

A1: Put the screenshot of Pintos running example here.

As the screenshot shows, the first booting is in QEMU and the second in Bochs.

```

  文件(F) 编辑(E) 查看(V) 转到(G) 运行(R) 终端(T) 帮助(H) 欢迎 - Pintos - Visual Studio Code
  问题 输出 调试控制台 终端
  PS D:\Pintos> docker run -it --rm --name pintos --mount type=bind,source=D:\Pintos,target=/home/PKU05/pintos,pkifyingng/pintos bash
  root@bac43aaaf41:~# ls
  pintos toolchain
  root@bac43aaaf41:~# cd pintos/src/threads
  root@bac43aaaf41:~/pintos/src/threads# cd build
  root@bac43aaaf41:~/pintos/src/threads/build# pintos --
  qemu-system-i386 -device isa-debug-exit -drive format=raw,media=disk,index=0,file=/tmp/_Nwo1NIJTJ.disk -m 4 -net none -nographic -monitor null
  Pintos ndat
  Loading.....
  Kernel command line:
  Pintos booting with 3,968 kB RAM...
  367 pages available in kernel pool.
  367 pages available in user pool.
  Calibrating timer... 185,472,000 loops/s.
  Boot complete.
  qemu-system-i386: terminating on signal 2

  root@bac43aaaf41:~/pintos/src/threads/build# pintos --bochs --
  squish-pty bochs -q
  =====
  Bochs x86 Emulator 2.6.2
  Built from SVN snapshot on May 26, 2013
  Compiled on Mar 1 2022 at 16:09:16
  =====
  000000000000[ ] reading configuration from bochsrc.txt
  000000000000[ ] bochsrc.txt:8: 'user_shortcut' will be replaced by new 'keyboard' option.
  000000000000[ ] installing ngui module as the Bochs GUI
  000000000000[ ] using log file bochsout.txt
  Pintos ndat
  Loading.....
  Kernel command line:
  Pintos booting with 4,696 kB RAM...
  383 pages available in kernel pool.
  383 pages available in user pool.
  Calibrating timer... 182,400 loops/s.
  Boot complete.
  
```

## Debugging

## QUESTIONS: BIOS

**B1: What is the first instruction that gets executed?**

`ljmp $0x3630,$0xf000e05b`. As shown in the screenshot.

```
(gdb) debugpintos
The target architecture is assumed to be i8086
[f000:fff0] 0xffff0: ljmp $0x3630,$0xf000e05b
0x0000fff0 in ?? ()
(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb) si
[f000:e062] 0xfe062: jne 0xd241d416
0x0000e062 in ?? ()
(gdb)
```

B2: At which physical address is this instruction located?

The PA is `0xffff0`.

As we know,  $PA = VA - PHYS\_BASE$ , and `PHYS_BASE` is a multiple of `0x10000000`. As the screenshot shows, VA of the first instruction is `0xf000fff0`. So we can easily induce that `PHYS_BASE = 0xf0000000`, `PA = 0xffff0`

Ref:

- <https://pkuflyingpig.gitbook.io/pintos/appendix/reference-guide/loading#physical-memory-map>
- <https://pkuflyingpig.gitbook.io/pintos/appendix/reference-guide/virtual-addresses#work-with-mapping-kernel-vm-one-to-one-to-pm>

## QUESTIONS: BOOTLOADER

B3: How does the bootloader read disk sectors? In particular, what BIOS interrupt is used?

First, the bootloader (beginning at `0x7c00`) checks all the sectors on different disks one by one.

```
46 00007c1e <read_mbr>:
47 > read_mbr: ...
77 > 00007c47 <check_partition>: ...
79 > check_partition: ...
103 00007c5e <next_partition>:
104
105 > next_partition: ...
114 > 00007c69 <next_drive>: ...
117 > next_drive: ...
124 00007c6d <no_boot_partition>:
125
```

Relevant functions.

When checking a specific sector, it uses `int 0x13` to call Interrupt 0x13 (Line 346, `0x7d30`). The number in register `%ah` is 0x42 (Line 342, `0x7d2c`). Looking up the table we know that the BIOS interrupt is "**Extended Read Sectors**".

```

322 read_sector:
323     pusha
324     7d1f: 60             pusha
325     sub %ax, %ax
326     7d20: 29 c0         sub    %eax,%eax
327     push %ax           # LBA sector number [48:63]
328     7d22: 50             push    %eax
329     push %ax           # LBA sector number [32:47]
330     7d23: 50             push    %eax
331     push %ebx          # LBA sector number [0:31]
332     7d24: 66 53         push    %bx
333     push %es           # Buffer segment
334     7d26: 06             push    %es
335     push %ax           # Buffer offset (always 0)
336     7d27: 50             push    %eax
337     push $1            # Number of sectors to read
338     7d28: 6a 01         push    $0x1
339     push $16           # Packet size
340     7d2a: 6a 10         push    $0x10
341     mov $0x42, %ah      # Extended read
342     7d2c: b4 42         mov     $0x42,%ah
343     mov %sp, %si        # DS:SI -> packet
344     7d2e: 89 e6         mov     %esp,%esi
345     int $0x13           # Error code in CF
346     7d30: cd 13         int     $0x13
347     popa               # Pop 16 bytes, preserve flags
348     7d32: 61             popa

```

	0Ch	Seek to Specified Track
	0Dh	Reset Fixed Disk Controller
	15h	Get Drive Type
	16h	Get Floppy Drive Media Change Status
	17h	Set Disk Type
	18h	Set Floppy Drive Media Type
	41h	Extended Disk Drive (EDD) Installation Check
	42h	Extended Read Sectors
	43h	Extended Write Sectors
	44h	Extended Verify Sectors

[https://en.wikipedia.org/wiki/BIOS\\_interrupt\\_call#Interrupt\\_table](https://en.wikipedia.org/wiki/BIOS_interrupt_call#Interrupt_table)

B4: How does the bootloader decides whether it successfully finds the Pintos kernel?

The bootloader checks if a partition meets three criteria:

- An used partition?
- A Pintos kernel partition?
- A bootable partition?

If yes, this is the bootable Pintos kernel. And we jump to `load_kernel` to load the kernel.

```

79 ~ check_partition:
80     # Is it an unused partition?
81     cml $0, %es:(%si)
82     7c47: 26 66 83 3c 00 74     cmpw   $0x74,%es:(%eax,%eax,1)
83     je next_partition
84     7c4d: 10 e8         adc     %ch,%al
85
86     # Print [1-4].
87     call putc
88     7c4f: ae             scas    %es:(%edi),%al
89     7c50: 00 26         add     %ah,(%esi)
90
91     # Is it a Pintos kernel partition?
92     cmpb $0x20, %es:4(%si)
93     7c52: 80 7c 04 20 75     cmpb   $0x75,0x20(%esp,%eax,1)
94     jne next_partition
95     7c57: 06             push    %es
96
97     # Is it a bootable partition?
98     cmpb $0x80, %es:(%si)
99     7c58: 26 80 3c 80 74     cmpb   $0x74,%es:(%eax,%eax,4)
100    je load_kernel
101    7c5d: 20             .byte 0x20

```

B5: What happens when the bootloader could not find the Pintos kernel?

If bootloader could not find the kernel, it calls *interrupt 18h* by the instruction `INT 0x18`. This interrupt outputs an error message telling BIOS that the bootloader fails to find the kernel.

```
no_such_drive:
no_boot_partition:
    # Didn't find a Pintos kernel partition anywhere, give up.
    call puts
7c6d: e8 77 00 0d 4e    call 4e0d7ce9 <__bss_start+0x4e0cfee9>
7c72: 6f               outsl %ds:(%esi),(%dx)
7c73: 74 20            je 7c95 <load_kernel+0x17>
7c75: 66 6f            outsw %ds:(%esi),(%dx)
7c77: 75 6e            jne 7ce7 <puts>
7c79: 64              fs
7c7a: 0d              .byte 0xd
7c7b: 00 cd            add %cl,%ch
    .string "\rNot found\r"

    # Notify BIOS that boot failed. See [IntrList].
    int $0x18
7c7d: 18              .byte 0x18
```

18h

Execute [Cassette BASIC](#): On IBM machines up to the early PS/2 line, this interrupt would start the ROM Cassette BASIC. Clones did not have this feature and different machines/BIOSes would perform a variety of different actions if INT 18h was executed, most commonly an error message stating that no bootable disk was present. Modern machines would attempt to [boot from a network](#) through this interrupt. On modern machines this interrupt will be treated by the BIOS as a signal from the bootloader that it failed to complete its task. The BIOS can then take appropriate next steps.<sup>[3]</sup>

[https://en.wikipedia.org/wiki/BIOS\\_interrupt\\_call#Interrupt\\_table](https://en.wikipedia.org/wiki/BIOS_interrupt_call#Interrupt_table)

B6: At what point and how exactly does the bootloader transfer control to the Pintos kernel?

`ljmp *start` (Line 225), the bootloader jumps to the start point of the kernel.

The bootloader saves the start address in position "start" (which is initially part of the loader's code).

```
212    mov $0x2000, %ax
213    7cbf: b8 00 20 8e c0    mov $0xc08e2000,%eax
214    mov %ax, %es
215    mov %es:0x18, %dx
216    7cc4: 26 8b 16          mov %es:(%esi),%edx
217    7cc7: 18 00             sbb %al,(%eax)
218    mov %dx, start
219    7cc9: 89 16             mov %edx,(%esi)
220    7ccb: d7               xlat %ds:(%ebx)
221    7ccc: 7c c7            j1 7c95 <load_kernel+0x17>
222    movw $0x2000, start + 2
223    7cce: 06              push %es
224    7ccf: d9 7c 00 20      fnstcw 0x20(%eax,%eax,1)
225    ljmp *start
226    7cd3: ff 2e           ljmp *(%esi)
227    7cd5: d7             xlat %ds:(%ebx)
228    7cd6: 7c             .byte 0x7c
```

## QUESTIONS: KERNEL

B7: At the entry of `pintos_init()`, what is the value of expression `init_page_dir[pd_no(ptov(0))]` in hexadecimal format?

The answer is 0.

Use gdb to set a breakpoint in `pintos_init`. And use `print` to print the value of expression.

```
(gdb) b pintos_init
Breakpoint 1 at 0xc00202b6: file ../../threads/init.c, line 82.
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0xc00202b6 <pintos_init>: push %ebp

Breakpoint 1, pintos_init () at ../../threads/init.c:82
(gdb) p init_page_dir[pd_no(ptov(0))]
=> 0xc000efef: int3
=> 0xc000efef: int3
$1 = 0
(gdb)
```

▼ to delete

`ptov(0)` transfers PA 0 to VA, which is  $VA = PA + PHYS\_BASE = 0xf0000000$ . `pd_no = VA >> PDSHIFT = VA >> (10+12) = d`

```
static inline void *ptov (uintptr_t paddr)
{
    ASSERT ((void *) paddr < PHYS_BASE);
    return (void *) (paddr + PHYS_BASE);
} // transfer PA to VA

static inline uintptr_t pd_no (const void *va) {
    return (uintptr_t) va >> PDSHIFT;
} // Obtains page directory index from a virtual address
```

B8: When `pallocc_get_page()` is called for the first time,

B8.1 what does the call stack look like?

Set a breakpoint in `pallocc_get_page` and run until the breakpoint.

Issue `bt` to see the call stack, which includes `start()`, `pintos_init()`, `paging_init()` and current `pallocc_get_page()`.

```
Breakpoint 1, pallocc_get_page (flags=(PAL_ASSERT | PAL_ZERO)) at ../../threads/pallocc.c:112
(gdb) bt
#0 pallocc_get_page (flags=(PAL_ASSERT | PAL_ZERO)) at ../../threads/pallocc.c:112
#1 0xc0020522 in paging_init () at ../../threads/init.c:219
#2 0xc002031b in pintos_init () at ../../threads/init.c:104
#3 0xc002013d in start () at ../../threads/start.S:180
```

B8.2 what is the return value in hexadecimal format?

Issue `finish` to run until the function finishes. The return value is `0xc0101000`, which is a void pointer.

```
(gdb) finish
Run till exit from #0 pallocc_get_page (flags=(PAL_ASSERT | PAL_ZERO)) at ../../threads/pallocc.c:112
=> 0xc0020522 <paging_init+17>: add $0x10,%esp
0xc0020522 in paging_init () at ../../threads/init.c:219
Value returned is $1 = (void *) 0xc0101000
```

B8.3 what is the value of expression `init_page_dir[pd_no(ptov(0))]` in hexadecimal format?

Use gdb `print`, and we see the value is `0`.

```
(gdb) p init_page_dir[pd_no(ptov(0))]
=> 0xc000ef7f: int3
=> 0xc000ef7f: int3
$2 = 0
```

B9: When `palloc_get_page()` is called for the third time,

B9.1 what does the call stack look like?

After the breakpoint is set, issue `c` three times to get to the third calling of `palloc_get_page()`.

Issue `bt` to see the call stack, which includes `start()`, `pintos_init()`, `thread_start()`, `thread_create()` and current `palloc_get_page()`.

```
(gdb) c
Continuing.
=> 0xc002328c <palloc_get_page>:      push    %ebp

Breakpoint 1, palloc_get_page (flags=PAL_ASSERT | PAL_ZERO) at ../../threads/palloc.c:112
(gdb) c
Continuing.
=> 0xc002328c <palloc_get_page>:      push    %ebp

Breakpoint 1, palloc_get_page (flags=PAL_ZERO) at ../../threads/palloc.c:112
(gdb) bt
#0  palloc_get_page (flags=PAL_ZERO) at ../../threads/palloc.c:112
#1  0xc0020bf9 in thread_create (name=0xc002ea51 "idle", priority=0, function=0xc0021028 <idle>, aux=0xc000efac) at ../../threads/thread.c:178
#2  0xc0020aee in thread_start () at ../../threads/thread.c:111
#3  0xc0020334 in pintos_init () at ../../threads/init.c:123
#4  0xc002013d in start () at ../../threads/start.S:180
```

B9.2 what is the return value in hexadecimal format?

Issue `fin`, and we see the return value is `0xc01030000`, which is a void pointer.

```
(gdb) fin
Run till exit from #0  palloc_get_page (flags=PAL_ZERO) at ../../threads/palloc.c:112
=> 0xc0020bf9 <thread_create+55>:      add     $0x10,%esp
0xc0020bf9 in thread_create (name=0xc002ea51 "idle", priority=0, function=0xc0021028 <idle>, aux=0xc000efac) at ../../threads/thread.c:178
Value returned is $3 = (void *) 0xc0103000
```

B9.3 what is the value of expression `init_page_dir[pd_no(ptov(0))]` in hexadecimal format?

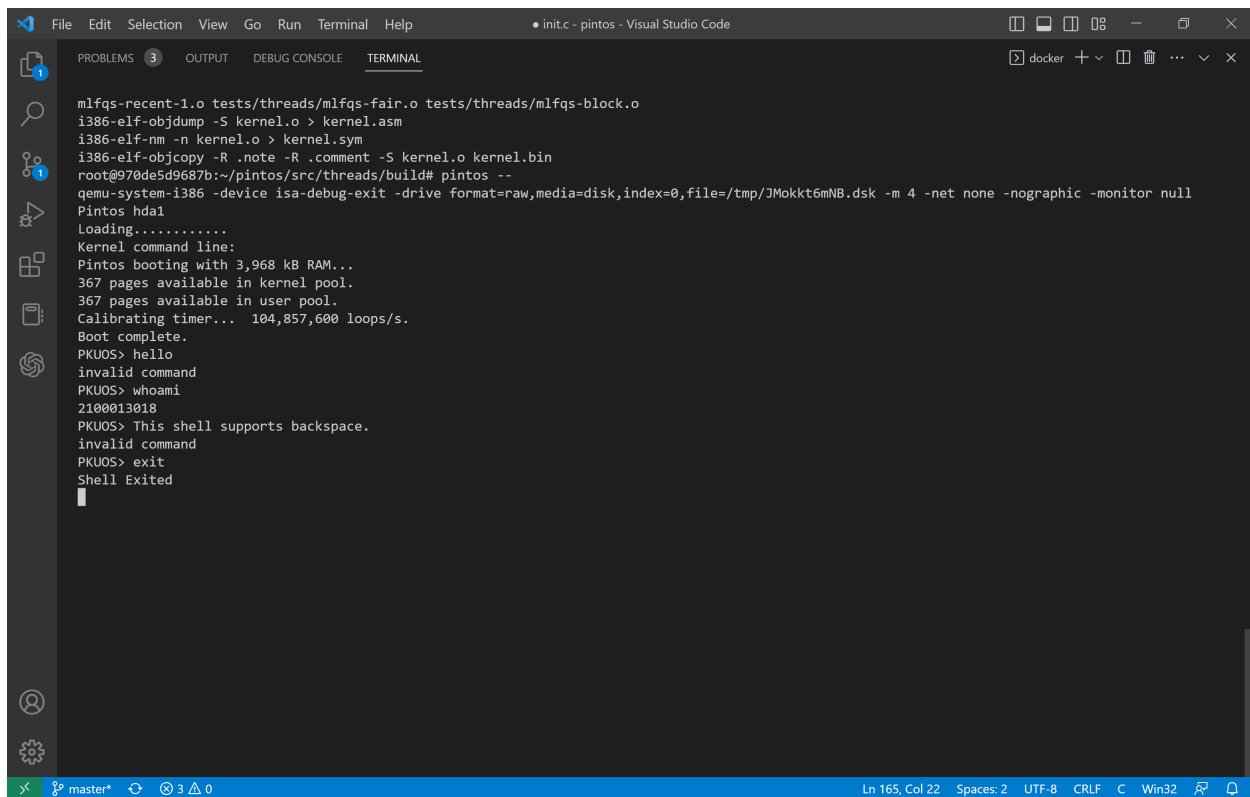
Use `gdb print` to get the value, which is `1056807` in decimal format.

Use `p/x` to get the hexadecimal value `0x102027`.

```
(gdb) p init_page_dir[pd_no(ptov(0))]
=> 0xc000ef3f: int3
=> 0xc000ef3f: int3
$4 = 1056807
(gdb) p/x init_page_dir[pd_no(ptov(0))]
=> 0xc000ef3f: int3
=> 0xc000ef3f: int3
$5 = 0x102027
```

## Kernel Monitor

C1: Put the screenshot of your kernel monitor running example here. (It should show how your kernel shell respond to `whoami`, `exit`, and other input.)



```
mlfq-recent-1.o tests/threads/mlfq-fair.o tests/threads/mlfq-block.o
i386-elf-objdump -S kernel.o > kernel.asm
i386-elf-nm -n kernel.o > kernel.sym
i386-elf-objcopy -R .note -R .comment -S kernel.o kernel.bin
root@970de5d9687b:~/pintos/src/threads/build# pintos --
qemu-system-i386 -device isa-debug-exit -drive format=raw,media=disk,index=0,file=/tmp/JMokkt6mNB.dsk -m 4 -net none -nographic -monitor null
Pintos hda1
Loading.....
Kernel command line:
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 104,857,600 loops/s.
Boot complete.
PKUOS> hello
invalid command
PKUOS> whoami
2100013018
PKUOS> This shell supports backspace.
invalid command
PKUOS> exit
Shell Exited
```

C2: Explain how you read and write to the console for the kernel monitor.

I use an infinite loop to print "PKUOS> " and get command line input. If a new line is detected, we enter the next loop.

```
144 while(print_prompt){
145     printf("%s", prompt);
146     int idx=0;
147     memset(cmdline, 0, MAXLINE);
148
149     // get cmdline
150 > while(1){...
172
173     // parse cmdline
174 > if(strcmp(cmdline, "whoami") == 0){...
177 > else if(strcmp(cmdline, "exit") == 0){...
180 > else{...
183 }
```

For each line, I write a `while` loop to get the chars one by one and:

- print it if printable
- delete a char if it is a backspace
- overlook it if the length of a line has exceeded `MAXLINE` ( a macro of max line length)
- break current loop and begin a new line if it is a `\r` (`\n` results in error)

```

149 // get cmdline
150 while(1){
151     char c = input_getc();
152     // newline dealing
153     if(c == '\n'){ // not \n
154         printf("\n");
155         cmdline[idx++] = '\0';
156         break;
157     }
158     // backspace dealing. Ref:https://blog.csdn.net/u011011827/article/details/125656962
159 > else if(c == 127){...
160     if(idx >= MAXLINE-1) continue; // do nothing
161     if(c >= 32) // printable
162         printf("%c", c);
163     cmdline[idx++] = c;
164 }
165 }

```

Note: As for the backspace, I refer to the article below.

```

158 // backspace dealing. Ref:https://blog.csdn
159 else if(c == 127){
160     if(idx > 0){
161         cmdline[idx--] = '\0';
162         // printf("\b"); // Errors Occur
163         printf("\b \b");
164     }
165     continue;
166 }

```

[https://blog.csdn.net/u011011827/article/details/125656962?ops\\_request\\_misc={\"request\\_id\": \"167772676616800188575949\", \"scm\": \"20140713.130102334.pc\\_all.\"}&request\\_id=167772676616800188575949&biz\\_task-code-2~all~first\\_rank\\_ecpm\\_v1~rank\\_v31\\_ecpm-1-125656962-0-null-null.142^v73^insert\\_down1,201^v4^add\\_ask,239^v2^insert\\_chatgpt&utm\\_term=shell中实现backspace](https://blog.csdn.net/u011011827/article/details/125656962?ops_request_misc={\)

(Note: We wrote a similar shell in ICS Shelllab, so I refer to my own code ( `ω´ ) )